



DCC

DEPARTAMENTO DE
CIÊNCIA DA COMPUTAÇÃO



UFMG

SPEC CPU 2017: Use cases

Andrei Rimsa Álvares

andrei@cefetmg.br



Summary

- Introduction
- Installing
- Configuring
- Validating
- Profiling
 - Dynamic profiling
 - Feedback-directed optimizations (FDO)

Introduction



Introduction

- SPEC[◊] (*Standard Performance Evaluation Corporation*) is a non-profit organization responsible for benchmarks that evaluates performance and energy efficiency of computer systems
 - Cloud
 - CPU
 - Graphics/Workstations
 - High Performance Computing
 - Java Client/Server
 - Machine Learning
 - Storage
 - Power
 - Virtualization

We are interested in SPEC CPU.



SPEC CPU

- SPEC CPU has a long history

- SPEC CPU 92
- SPEC CPU 95
- SPEC CPU 2000
- SPEC CPU 2006
- SPEC CPU 2017[‡]

We are interested in the latest release: SPEC CPU 2017.

Info: a license is required to use SPEC CPU 2017.

[‡]: <https://www.spec.org/cpu2017/>



SPEC CPU 2017

- SPEC CPU 2017 contains a set of CPU intensive suites for measuring and comparing compute intensive performance, stressing a system's processor, memory subsystem and compiler
 - It is package with 43 benchmarks organized into four suites
 - SPECspeed® 2017 Integer
 - SPECspeed® 2017 Floating Point
 - SPECCrate® 2017 Integer
 - SPECCrate® 2017 Floating Point
-
-



SPEC Suites

- Overview of SPEC suites[◊]

base: must use the same compiler flags.
peak: may use different compiler flags.

Short Tag	Suite	Contents	Metrics	How many copies? What do Higher Scores Mean?
intspeed	SPECspeed® 2017 Integer	10 integer benchmarks	SPECspeed@2017_int_base SPECspeed@2017_int_peak SPECspeed@2017_int_energy_base SPECspeed@2017_int_energy_peak	SPECspeed suites always run one copy of each benchmark. Higher scores indicate that less time is needed.
fpspeed	SPECspeed®2017 Floating Point	10 floating point benchmarks	SPECspeed@2017_fp_base SPECspeed@2017_fp_peak SPECspeed@2017_fp_energy_base SPECspeed@2017_fp_energy_peak	
intrate	SPECrate® 2017 Integer	10 integer benchmarks	SPECrate@2017_int_base SPECrate@2017_int_peak SPECrate@2017_int_energy_base SPECrate@2017_int_energy_peak	SPECrate suites run multiple concurrent copies of each benchmark. The tester selects how many. Higher scores indicate more <i>throughput</i> (work per unit of time).
fprate	SPECrate® 2017 Floating Point	13 floating point benchmarks	SPECrate@2017_fp_base SPECrate@2017_fp_peak SPECrate@2017_fp_energy_base SPECrate@2017_fp_energy_peak	

The "Short Tag" is the canonical abbreviation for use with `runcpu`, where context is defined by the tools. In a published document, context may not be clear. To avoid ambiguity in published documents, the Suite Name or the Metrics should be spelled as shown above.

We will reference some collection of benchmarks by their short name.

[◊]: <https://www.spec.org/cpu2017/Docs/overview.html#suites>

SPECinteger

- Collection of SPEC Integer benchmarks

SPECrate®2017 Integer	SPECSspeed®2017 Integer	Language [1]	KLOC [2]	Application Area
500.perlbench_r	600.perlbench_s	C	362	Perl interpreter
502.gcc_r	602.gcc_s	C	1,304	GNU C compiler
505.mcf_r	605.mcf_s	C	3	Route planning
520.omnetpp_r	620.omnetpp_s	C++	134	Discrete Event simulation - computer network
523.xalancbmk_r	623.xalancbmk_s	C++	520	XML to HTML conversion via XSLT
525.x264_r	625.x264_s	C	96	Video compression
531.deepsjeng_r	631.deepsjeng_s	C++	10	Artificial Intelligence: alpha-beta tree search (Chess)
541.leela_r	641.leela_s	C++	21	Artificial Intelligence: Monte Carlo tree search (Go)
548.exchange2_r	648.exchange2_s	Fortran	1	Artificial Intelligence: recursive solution generator (Sudoku)
557.xz_r	657.xz_s	C	33	General data compression

SPECrate

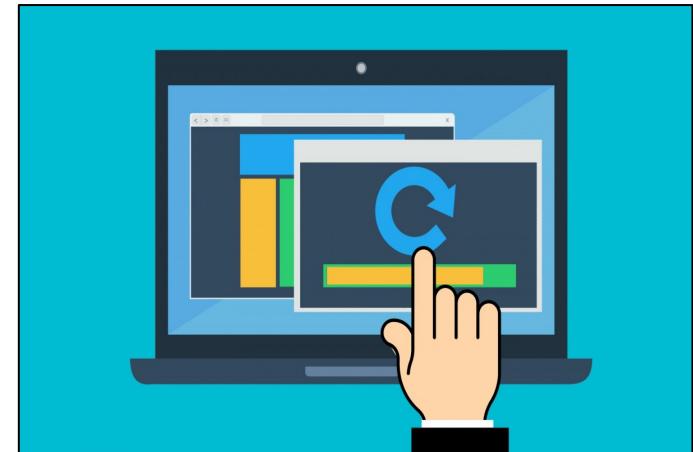
- Collection of SPEC Floating Point benchmarks

SPECrate®2017 Floating Point	SPECspeed®2017 Floating Point	Language [1]	KLOC [2]	Application Area
503.bwaves_r	603.bwaves_s	Fortran	1	Explosion modeling
507.cactuBSSN_r	607.cactuBSSN_s	C++, C, Fortran	257	Physics: relativity
508.namd_r		C++	8	Molecular dynamics
510.parest_r		C++	427	Biomedical imaging: optical tomography with finite elements
511.povray_r		C++, C	170	Ray tracing
519.lbm_r	619.lbm_s	C	1	Fluid dynamics
521.wrf_r	621.wrf_s	Fortran, C	991	Weather forecasting
526.blender_r		C++, C	1,577	3D rendering and animation
527.cam4_r	627.cam4_s	Fortran, C	407	Atmosphere modeling
	628.pop2_s	Fortran, C	338	Wide-scale ocean modeling (climate level)
538.imagick_r	638.imagick_s	C	259	Image manipulation
544.nab_r	644.nab_s	C	24	Molecular dynamics
549.fotonik3d_r	649.fotonik3d_s	Fortran	14	Computational Electromagnetics
554.roms_r	654.roms_s	Fortran	210	Regional ocean modeling

[1] For multi-language benchmarks, the first one listed determines library and link options ([details](#))

[2] KLOC = line count (including comments/whitespace) for source files used in a build / 1000

Installing



Prerequisites

- Ensure we have compilers for C, C++ and Fortran

- For *gcc* and *g++*

```
$ sudo apt-get install build-essential  
...  
$ gcc --version  
gcc (Ubuntu 9.3.0-17ubuntu1~20.04) 9.3.0  
...  
$ g++ --version  
g++ (Ubuntu 9.3.0-17ubuntu1~20.04) 9.3.0  
...
```

- For *gfortran*

```
$ sudo apt-get install gfortran  
...  
$ gfortran --version  
GNU Fortran (Ubuntu 9.3.0-17ubuntu1~20.04) 9.3.0  
...
```

Installing

- Install SPEC CPU 2017 in three steps

1) Mount the iso image

- With root privileges (using *mount*)

```
$ mkdir /tmp/iso  
$ sudo mount -o loop,ro cpu2017-1_0_5.iso /tmp/iso
```

- Without root privileges (using *fuseiso*)

```
$ mkdir /tmp/iso  
$ fuseiso /tmp/iso cpu2017-1_0_5.iso
```

Installing

- Install SPEC CPU 2017 in three steps
 - 2) Call the *install.sh* script and select the installation directory

```
$ cd /tmp/iso  
$ mkdir ~/benchmarks  
$ ./install.sh  
SPEC CPU2017 Installation
```

```
Top of the CPU2017 tree is '/tmp/iso'  
Enter the directory you wish to install to (e.g. /usr/cpu2017)  
/home/andrei/benchmarks/SPEC17
```

```
Installing FROM /tmp/iso  
Installing TO /home/andrei/benchmarks/SPEC17
```

```
Is this correct? (Please enter 'yes' or 'no')  
yes  
...
```

Installing

- Install SPEC CPU 2017 in three steps

3) Unmount the iso image

- With root privileges (using *umount*)

```
$ cd -
$ sudo umount /tmp/iso
$ rmdir /tmp/iso
```

- Without root privileges (using *fusermount*)

```
$ cd -
$ fusermount -u /tmp/iso
$ rmdir /tmp/iso
```

Directory Structure

- The structure of the CPU 2017 directory tree is
 - **\$SPEC or %SPEC%:** the root directory
 - **benchspec:** some suite-wide files
 - **CPU:** the benchmarks
 - **bin:** tools to run and report on the suite
 - **config:** config files
 - **Docs:** HTML and plaintext documentation
 - **result:** log files and reports
 - **tmp:** temporary files
 - **tools:** sources for the CPU 2017 tools

Directory Structure

- Within each of the individual benchmarks, the structure is
 - **nnn.benchmark**: root for this benchmark
 - **build**: benchmark binaries are built here
 - **data**
 - **all**: data used by all runs (if needed by the benchmark)
 - **ref**: the timed data set
 - **test**: data for a simple test that an executable is functional
 - **train**: data for feedback-directed optimization
 - **Docs**: documentation for this benchmark
 - **exe**: compiled versions of the benchmark
 - **run**: benchmarks are run here
 - **Spec**: SPEC metadata about the benchmark
 - **src**: The sources for the benchmark

Configuring



Configuring

- Configure SPEC CPU 17 in three steps

1) Copy base template files

```
$ cd ~/benchmarks/SPEC17  
$ cp config/Example-gcc-linux-x86.cfg config/andrei.cfg  
$ cp config/flags/gcc.xml config/flags/andrei-gcc.xml
```

Unsupported compiler flags can be included for reportable runs.

Enter SPEC CPU 2017 directory.

Copy a base template file such as gcc for x86.

Info: some configuration options can be overwritten at runtime.

Configuring

- Configure SPEC CPU 17 in three steps
 - Add a flag (`-no-pie`) to the flags file (*andrei-gcc.xml*)

```
<flag name="F-mvis3"
      class="optimization">
  <![CDATA[<p>
    Generate code to take advantage of version 3 of the SPARC Visual Instruction Set extensions
  </p>]]>
</flag>

<flag name="F-no-pie"
      class="optimization">
  <![CDATA[<p>
    Do not generate position independent executable (PIE).
  </p>]]>
</flag>

<flag name="F-Ofast"
      class="optimization">
  <![CDATA[<p>
    Enable all optimizations of -O3 plus optimizations that are not valid for standard-compliant
    programs, such as re-ordering operations without regard to parentheses.
    <br /><a href="#gnote">Many more details are available</a>.
  </p>]]>
</flag>
```



New entry for flag `-no-pie`.

Configuring

- Configure SPEC CPU 17 in three steps
 - 3) Edit the config file (*andrei.cfg*)

```
...
#----- Global Settings -----
...
flagsurl      = ${top}/config/flags/andrei-gcc.xml
...
tune          = base
...
#----- How Many CPUs? -----
...
intrate,fprate:
  copies      = 1
intspeed,fpspeed:
  threads     = 4
...
#----- Compilers -----
...
#ifndef %{gcc_dir}
#define gcc_dir      /usr
#endif
...
#----- Baseline Tuning Flags -----
...
default=base:
  OPTIMIZE     = -g -O3 -march=native -no-pie -fno-unsafe-math-optimizations -fno-tree-loop-vectorize
...
```

Keep as *base* for now.

Use the new flags file.

One instance for *intrate* and *fprate*.

Four threads for *intspeed* and *fpspeed*.

Path to *gcc/g++/gfortran*.

Add the *-no-pie* optimization flag for *base*.

Validating



Load simulation environment.

Validating

- Use runcpu[♡] to run SPEC CPU 2017

Can use sets (ex.: *inrate, fprate*), benchmarks (ex.: *502.gcc_r*) or supersets (ex.: *specrate, specspeed* or *all*).

Reportable runs are suitable for public reporting

```
$ . shrc
$ runcpu --config=andrei.cfg --iterations=3 --reportable inrate
SPEC CPU(r) 2017 Benchmark Suites
Copyright 1995-2017 Standard Performance Evaluation Corporation (SPEC)
runcpu v5825
Using 'linux-x86_64' tools
Reading file manifests... read 32272 entries from 2 files in 0.22s (145963 files/s)
Loading runcpu modules.....
Locating benchmarks...found 47 benchmarks in 53 benchsets.
Reading config file '/home/andrei/benchmarks/SPEC17/config/andrei.cfg'
1 configuration selected:
```

Action	Run Mode	Workload	Report Type
validate	rate	refrate	SPECrate2017_int

Valid runs must execute each benchmark at least three times.

Warning: ulimit -s unlimited may be required for 627.cam4_s (*fpspeed*)

Setting up environment for running inrate...
Starting runcpu for inrate...
...

The default action is *validate*.

[♡]: <https://www.spec.org/cpu2017/Docs/runcpu.html>

Valid Results

- For reportable runs, a results report should look something like this

SPEC CPU2017 Integer Rate Result												
Copyright 2017-2021 Standard Performance Evaluation Corporation												
My Corporation								SPECRate2017_int_base = 2.62				
CPU2017 License: nnn (Your SPEC license number)								Test Date: Dec-2021				
Test Sponsor: My Corporation								Hardware Availability:				
Tested by: My Corporation								Software Availability:				
Results Table												
Benchmark	Copies	Seconds	Ratio	Seconds	Ratio	Seconds	Ratio	Copies	Seconds	Ratio	Seconds	Ratio
500.perlbench_r	1	619	2.57	618	2.57	617	2.58					
502.gcc_r	1	413	3.43	412	3.43	412	3.43					
505.mcf_r	1	523	3.09	522	3.09	521	3.10					
520.omnetpp_r	1	608	2.16	602	2.18	605	2.17					
523.xalancbmk_r	1	618	1.71	600	1.76	607	1.74					
525.x264_r	1	657	2.66	659	2.66	659	2.66					
531.deepsjeng_r	1	476	2.41	476	2.41	475	2.41					
541.leela_r	1	733	2.26	735	2.25	733	2.26					
548.exchange2_r	1	487	5.38	483	5.42	483	5.42					
557.xz_r	1	563	1.92	563	1.92	564	1.92					
SPECRate2017_int_base = 2.62												
SPECRate2017_int_peak = Not Run												
Results appear in the order in which they were run. Bold underlined text indicates a median measurement.												
General Notes												
Environment variables set by runcpu before the start of the run: LD_LIBRARY_PATH = "/usr/lib64:/usr/lib:/lib64"												
Platform Notes												
Sysinfo program /home/andrei/benchmarks/SPEC17/bin/sysinfo Rev: r5974 of 2018-05-19 9bcde8f2999c33d61f64985e45859ea9 running on opencl Fri Dec 3 10:14:11 2021												
SUT (System Under Test) info as seen by some common utilities. For more information on this section, see https://www.spec.org/cpu2017/Docs/config.html#sysinfo												
From /proc/cpuinfo model name : Intel(R) Xeon(R) CPU E5-2620 0 @ 2.00GHz 1 "physical id"s (chips) 12 "processors" cores, siblings (Caution: counting these is hw and system dependent. The following excerpts from /proc/cpuinfo might not be reliable. Use with caution.) cpu cores : 6 siblings : 12 physical 0: cores 0 1 2 3 4 5												
From lscpu: Architecture: x86_64 CPU op-mode(s): 32-bit, 64-bit												
(Continued on next page)												
Page 2	Standard Performance Evaluation Corporation (info@spec.org)								https://www.spec.org/			

SPEC CPU2017 Integer Rate Result

Copyright 2017-2021 Standard Performance Evaluation Corporation

Invalid Results

- Otherwise, in case of any problems, an "Invalid Run" warning will be displayed in the report

Info: if the problems were caused by flags, the report can be recovered with rawformat[♪].

My Corporation		SPECrate2017_int_base = 2.62																	
SPECrate2017_int_peak = Not Run																			
CPU2017 License: nnn (Your SPEC license number)																			
Test Sponsor: My Corporation																			
Tested by: My Corporation																			
Results Table																			
Benchmark	Base						Peak												
	Copies	Seconds	Ratio	Seconds	Ratio	Seconds	Ratio	Copies	Seconds	Ratio	Seconds	Ratio							
500.perlbench_r	1	619	2.57	618	2.57	617	2.58												
502.gcc_r	1	413	3.43	412	3.43	412	3.43												
505.mcf_r	1	523	3.09	522	3.09	521	3.10												
520.omnetpp_r	1	608	2.16	602	2.18	605	2.17												
523.xalancbmk_r	1	618	1.71	600	1.76	607	1.74												
525.x264_r	1	657	2.66	659	2.66	659	2.66												
531.deepsjeng_r	1	476	2.41	476	2.41	475	2.41												
541.leela_r	1	733	2.26	735	2.25	733	2.26												
548.exchange2_r	1	487	5.39	483	5.42	483	5.42												
557.xz_r	1	563	1.92	563	1.92	564	1.92												
SPECrate2017_int_base = 2.62				SPECrate2017_int_peak = Not Run															
Results appear in the order in which they were run. Bold underlined text indicates a median measurement.																			
General Notes																			
Environment variables set by runcpu before the start of the run: LD_LIBRARY_PATH = "/usr/lib64/:/usr/lib/:/lib64"																			
Platform Notes																			
Sysinfo program /home/andrei/benchmarks/SPEC17/bin/sysinfo Rev: r5974 on 2018-05-19 9bcde8f2999c33d61f64985e45859ea9 running on opencl Fri Dec 3 10:14:11 2021																			
SUT (System Under Test) info as seen by some common utilities. For more information on this section, see https://www.spec.org/cpu2017/Docs/config.html#sysinfo																			
From /proc/cpuinfo model name : Intel(R) Xeon(R) CPU E5-2620 0 @ 2.00GHz 1 "physical id"s (chips) 12 "processors" cores, siblings (Caution: counting these is hw and system dependent. The following excerpts from /proc/cpuinfo might not be reliable. Use with caution.) cpu cores : 6 siblings : 12 physical 0: cores 0 1 2 3 4 5																			
From lscpu: Architecture: x86_64 CPU op-mode(s): 32-bit, 64-bit																			
(Continued on next page)																			
Page 2			Standard Performance Evaluation Corporation (info@spec.org)						https://www.spec.org/										

Scripting the Results

- The results is generated in different formats, including a textual one that can be scripted and output in a CSV-like format

```
$ (echo "benchmark,run1,run2,run3"; \  
  sed -n '/^---/,/^==/p' result/CPU2017.001.intrate.txt | \  
  awk '$1 ~ /^[[:digit:]]{3}/ {bench[$1]=bench[$1] " " $3} \  
      END { for (b in bench) print b bench[b]; }' | \  
  tr " " "," | sort -n)  
benchmark,run1,run2,run3  
600.perlbench_s,615,617,616  
602.gcc_s,780,778,783  
605.mcf_s,1059,1061,1066  
620.omnetpp_s,611,612,614  
623.xalancbmk_s,588,602,588  
625.x264_s,658,668,667  
631.deepsjeng_s,579,580,579  
641.leela_s,731,732,732  
648.exchange2_s,485,487,513  
657.xz_s,1281,1274,1276
```

Profiling



Profiling

- In the context of SPEC CPU 2017, we are interested in
 - **Dynamic profiling**: simulate a program monitoring its execution
 - Can use a Dynamic Binary Instrumentation (DBI) framework
 - Ex.: valgrind, pin, DynamoRIO, etc.
 - **Feedback-directed optimizations (FDO)**: compiled in two phases
 - In the first phase, the compiler adds instrumentation to monitor a program execution
 - Ex.: gcc -fprofile-generate ...
 - In the second phase, the compiler is feed with information collected at runtime
 - Ex.: gcc -fprofile-use ...

Dynamic Profiling

- To perform dynamic profiling, we need to prepend a DBI we intend to use to the command that executes the benchmark
 - Ex.: `valgrind -q --tool=cfggrind♡ -- benchmark args`
- However, there is no native way of doing that, thus a patch[◊] was created to support such feature

```
$ cd ~/benchmarks/SPEC17
$ wget http://lac.dcc.ufmg.br/misc/spec2017-prepend.patch
$ patch -p1 < spec2017-prepend.patch
patching file MANIFEST
patching file bin/harness/benchmark.pm
patching file bin/harness/configpp
patching file bin/harness/parse.pl
patching file bin/harness/runcpu
```

- And now we can use the `--prepend` option

```
$ . shrc
$ runcpu --config=andrei.cfg --iterations=3 --reportable \
--prepend="valgrind -q --tool=cfggrind --" intrate
```

[♡]: <https://github.com/rimsa/CFGgrind>

[◊]: <http://lac.dcc.ufmg.br/misc/SPEC2017.html>

Feedback-Directed Optimizations

- Let's use BOLT[♀] to optimize a program with the following workflow

```
$ wget https://raw.githubusercontent.com/rimsa/CFGgrind/master/tests/test.c
$ gcc -g -O3 -no-pie -o test test.c -Wl,-q
$ for r in {1..3}; do \
    numbers=( $(seq 1 10000 | sort -R) ); \
    perf record -e cycles:u -j any,u -o run${r}.data -- ./test ${numbers[@]}; \
    perf2bolt -p run${r}.data -o run${r}.fdata ./test; \
done
$ merge-fdata run{1..3}.fdata > all.fdata
$ llvm-bolt ./test -o ./test.opt -report-stale -data=all.fdata \
    -reorder-blocks=cache+ -reorder-functions=hfsort -split-functions=2 \
    -split-all-cold -split-eh -dyno-stats
```

How can we apply
this workflow to
SPEC CPU 2017?

[♀]: <https://github.com/facebookincubator/BOLT>

Configuring

- Add `-Wl,-q` to the flags file (*andrei-gcc.xml*)

```
...
</flag>

<flag name="Wl-q"
      class="optimization"
      regexp="-Wl,-q">
    <example>-Wl,-q</example>
    <! [CDATA[<p>
      Add the linker flag that enables relocation.
    </p>]]>
</flag>

<flag name="Wl-rpath"
      class="optimization"
      regexp="-Wl,-rpath,(\S+)(?:\s|$)">
    >
    <example>-Wl,-rpath,/path/to/lib</example>
    Add the specified directory to the runtime library search path used
    when linking an ELF executable with shared objects.
</flag>

...
```



New entry for flag `-Wl,-q`.

Configuring

- Update the config file (*andrei.cfg*)

```
...
#----- Global Settings -----
...
flagsurl      = ${top}/config/flags/andrei-gcc.xml
...
tune          = peak
...
#----- Peak Tuning Flags -----
default=peak:
# basepeak = yes # if you develop some peak tuning, remove this line.
OPTIMIZE = -g -O3 -march=native -no-pie -fno-unsafe-math-optimizations -fno-tree-loop-vectorize -Wl,-q
intrate,intspeed=peak: # flags for integer peak
EXTRA_COPTIMIZE = -fno-strict-aliasing -fgnu89-inline
...

```

Change *tune* to *peak* (required for FDO).

Comment or remove this line.

Copy and edit this section from base to peak.

And include *-Wl,-q* required for BOLT.

Copy *OPTIMIZE* option from *base* to *peak*.

Feedback-Directed Optimizations

- SPEC CPU 2017 has builtin support for Feedback-directed optimizations[‡] typically using the following steps
 - Build a program with instrumentation
 - Run sample training workload to collect profile information
 - Apply the profile to produce a modified program
- This behavior can be controlled by
 - PASSn makevars
 - fdo shell options
 - *feedback config file* option
 - --feedback runcpu option



We will only see how to do FDO with this option.

[‡]: <https://www.spec.org/cpu2017/Docs/config.html#UsingFeedback>

FDO Shell Options

- Some useful FDO shell options[‡]

- **fdo_pre0**: Commands to execute before starting an FDO series
- **fdo_preN**: Commands to be executed before pass N
- **fdo_pre_makeN**: Commands to be done prior to Nth compile
- **fdo_make_passN**: Commands to actually do the Nth compile
- **fdo_post_makeN**: Commands to be done after the Nth compile
- **fdo_runN**: Commands to be used for Nth training run
- **fdo_postN**: Commands to be done at the end of pass N

Info: we can write shell scripts in these options.

We will work only with these options.

[‡]: <https://www.spec.org/cpu2017/Docs/config.html#shellOptions>

Curly Variables

- Some useful curly variables[‡]
 - **`${baseexe}`**: The first part of the executable name, which is `<baseexe>_<tune>.<label>`. For example, in "lmb_r_base.foo", baseexe is "lmb_r".
 - **`${benchmark}`**: The number and name of the benchmark currently being run, for example 519.lmb_r
 - **`${benchname}`**: The name of the benchmark currently being run, for example lmb_r
 - **`${benchnum}`**: The number of the benchmark currently being run, for example 519
 - **`${command}`**: The shell command line to run the current benchmark, for example
`./run_base_test_sticky.0000/lmb_r_base.sticky 20 reference.dat 0 1
100_100_130_cf_a.of`
 - **`${commandexe}`**: The executable for the current command, for example
`./run_base_test_none.0000/lmb_r_base.sticky`
 - **`${iter}`**: The current iteration number
 - **`${tune}`**: The tuning for the benchmark being run (base or peak)
 - **`${workload}`**: The current workload number (within the iteration)N

[‡]: <https://www.spec.org/cpu2017/Docs/config.html#sectionI.F.5>

Configuring

- Update the config file (*andrei.cfg*) to include FDO

```
#----- Peak Tuning Flags -----  
default=peak:  
    basepeak = yes # if you develop some peak tuning, remove this line.  
  
fdo_pre0 = rm -rf ${benchmark}.data ${benchmark}.fdata; \  
           touch ${benchmark}.fdata  
  
fdo_run1 = mv ${benchmark}.fdata previous.fdata; \  
           perf record -e cycles:u -j any,u -o ${benchmark}.data -- ${command}; \  
           perf2bolt -p ${benchmark}.data -o current.fdata ${baseexe}; \  
           merge-fdata previous.fdata current.fdata > ${benchmark}.fdata; \  
           rm -rf ${benchmark}.data previous.fdata current.fdata  
  
fdo_post1 = if [ -s "${benchmark}.fdata" ]; then \  
           mv ${baseexe} ${baseexe}-orig; \  
           llvm-bolt ${baseexe}-orig -o ${baseexe} -report-stale -data=${benchmark}.fdata  
-reorder-blocks=cache+ -reorder-functions=hfsort -split-functions=2 -split-all-cold -split-eh  
-dyno-stats; \  
           fi  
...
```

Clean up previous runs and create empty fdata using \${benchmark}.

Profile \${command} with perf and convert output to fdata.

Optimize \${baseexe} with llvm-bolt.

Debugging

- Debug FDO using --fake.

```
$ runcpu --config=andrei.cfg --fake 502.gcc_r
...
%% Fake commands from fdo_pre0 (rm -rf ${benchmark}.data ${benchmark}...):
rm -rf 502.gcc_r.data 502.gcc_r.fdata; touch 502.gcc_r.fdata
%% End of fake output from fdo_pre0 (rm -rf ${benchmark}.data ${benchmark}...)
...
# Starting run for copy #0
cd /home/andrei/benchmarks/SPEC17/benchspec/CPU/502.gcc_r/build/build_peak_mytest-m64.0000
mv 502.gcc_r.fdata previous.fdata; perf record -e cycles:u -j any,u -o 502.gcc_r.data --
..../build_peak_mytest-m64.0000/cpugcc_r 200.c -O3 -finline-limit=50000 -o 200.opts-O3-
-finline-limit_50000.s > 200.opts-O3-finline-limit_50000.out 2>> 200.opts-O3-finline-
limit_50000.err; perf2bolt -p 502.gcc_r.data -o current.fdata cpugcc_r; merge-fdata
previous.fdata current.fdata > 502.gcc_r.fdata; rm -rf 502.gcc_r.data previous.fdata
current.fdata
...
%% End of fake output from benchmark_run (/home/andrei/benchmarks/SPEC17/bin/speci...)
...
%% Fake commands from fdo_post1 (if [ -s "${benchmark}.fdata" ]; then...):
if [ -s "502.gcc_r.fdata" ]; then mv cpugcc_r cpugcc_r-orig; llvm-bolt cpugcc_r-orig -o
cpugcc_r -report-stale -data=502.gcc_r.fdata -reorder-blocks=cache+ -reorder-
functions=hfsort -split-functions=2 -split-all-cold -split-eh -dyno-stats; fi
%% End of fake output from fdo_post1 (if [ -s "${benchmark}.fdata" ]; then...)
...
```

The diagram illustrates the flow of fake commands generated by the runcpu tool. It highlights three distinct phases: **fdo_pre0**, **fdo_run1**, and **fdo_post1**. Each phase is represented by a green callout box with a bracket connecting it to its corresponding command block in the log output. The **fdo_pre0** box covers the initial cleanup and touch operations. The **fdo_run1** box covers the main benchmark execution and data merging. The **fdo_post1** box covers the final cleanup and bolting of the benchmark data.

Executing

- If everything is alright, execute the benchmark

```
$ runcpu --config=andrei.cfg \
--iterations=3 --reportable \
intrate
```

SPEC CPU2017 Integer Rate Result														
Copyright 2017-2021 Standard Performance Evaluation Corporation														
My Corporation														SPECrate2017_int_base = 2.62
														SPECrate2017_int_peak = 2.67
CPU2017 License: nnn (Your SPEC license number)														Test Date: Dec-2021
Test Sponsor: My Corporation														Hardware Availability:
Tested by: My Corporation														Software Availability:
Results Table														
Benchmark	Base							Peak						
	Copies	Seconds	Ratio	Seconds	Ratio	Seconds	Ratio	Copies	Seconds	Ratio	Seconds	Ratio	Seconds	Ratio
500.perlbench_r	1	616	2.59	616	2.59	618	2.58	1	576	2.76	578	2.76	574	2.77
502.gcc_r	1	412	3.43	412	3.44	410	3.45	1	389	3.64	388	3.65	389	3.64
505.mcf_r	1	523	3.09	523	3.09	520	3.11	1	524	3.08	525	3.08	523	3.09
520.omnetpp_r	1	612	2.14	607	2.16	604	2.17	1	598	2.19	601	2.18	598	2.19
523.xalancbmk_r	1	608	1.74	610	1.73	605	1.75	1	596	1.77	604	1.75	601	1.76
525.x264_r	1	659	2.66	659	2.66	659	2.66	1	649	2.70	650	2.69	651	2.69
531.deepsjeng_r	1	475	2.41	475	2.41	474	2.42	1	471	2.43	471	2.43	471	2.43
541.leela_r	1	732	2.26	733	2.26	733	2.26	1	734	2.26	732	2.26	732	2.26
548.exchange2_r	1	490	5.35	489	5.36	485	5.41	1	492	5.32	491	5.33	490	5.35
557.xz_r	1	562	1.92	563	1.92	563	1.92	1	558	1.94	561	1.92	558	1.93
SPECrate2017_int_base = 2.62														SPECrate2017_int_peak = 2.67
Results appear in the order in which they were run. Bold underlined text indicates a median measurement.														
General Notes														
Environment variables set by runcpu before the start of the run: LD_LIBRARY_PATH = "/usr/lib64:/usr/lib:/lib64"														
Platform Notes														
Sysinfo program /home/andrei/benchmarks/SPEC17/bin/sysinfo Rev: r5974 of 2018-05-19 9bcde8f2999c33d3d61f64985e45859ea9 running on opencl Thu Dec 2 22:59:43 2021														
SUT (System Under Test) info as seen by some common utilities. For more information on this section, see https://www.spec.org/cpu2017/Docs/config.html#sysinfo														
From /proc/cpuinfo model name : Intel(R) Xeon(R) CPU E5-2620 0 @ 2.00GHz 1 "physical id"s (chips) 12 "processors" cores, siblings (Caution: counting these is hw and system dependent. The following excerpts from /proc/cpuinfo might not be reliable. Use with caution.) cpu cores : 6 siblings : 12 physical 0: cores 0 1 2 3 4 5														
From lscpu: Architecture: x86_64 CPU op-mode(s): 32-bit, 64-bit														
(Continued on next page)														

Questions?

